

Table 1: Zero-shot + cache-scaffold only (all prompts include [CACHE]; not the plain eval_babilongv2 baseline). First row: scaffold-only control (metrics/scaffold_baselines/scaffold_loss_k4_layer62.json, baseline_acc). Other rows: cache_acc matching ablation_results_table.tex; sources metrics/{run_id}.json, except layer sweep rows (Loss-selected L6/L16, Uniform interval L2) from metrics/layer-experiments/{run_id}.json. QA1-QA3, 50 per cell, inline eval.

Method	0k	1k	2k	4k	8k	16k	Avg.
Scaffold only, no StateCache ($k=4$ loss, L62)	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Loss-selected, $k = 1$, L62	22.7	11.3	20.0	16.0	20.0	14.7	17.4
Loss-selected, $k = 4$, L62	16.7	18.0	15.3	12.7	14.7	14.7	15.3
Loss-selected, $k = 8$, L62	20.0	18.7	15.3	21.3	16.7	22.0	19.0
Loss-selected, $k = 16$, L62	18.0	12.7	22.7	11.3	20.0	13.3	16.3
Loss-selected, $k = 32$, L62	23.3	10.0	19.3	12.7	18.7	19.3	17.2
Loss-selected, $k = 64$, L62	15.3	18.7	25.3	18.0	14.0	21.3	18.8
Loss-selected, $k = 4$, L62, train up to 8k	20.0	14.7	16.7	15.3	13.3	15.3	15.9
Loss-selected, $k = 4$, L62, no gate	18.0	15.3	14.0	14.7	10.7	16.7	14.9
Loss-selected, $k = 4$, L62, full loss	21.3	18.7	20.7	9.3	16.7	19.3	17.7
Loss-selected, $k = 4$, L32	48.7	44.7	33.3	31.3	25.3	19.3	33.8
Loss-selected, $k = 4$, L2	62.0	56.7	48.0	35.3	38.0	24.7	44.1
Loss-selected, $k = 4$, L6	65.3	60.7	49.3	48.0	37.3	26.7	47.9
Loss-selected, $k = 4$, L16	64.7	53.3	42.0	33.3	34.7	24.7	42.1
Uniform interval, $k = 4$, L2	63.3	54.7	42.0	37.3	34.0	21.3	42.1
Uniform interval, $k = 4$, L62	21.3	20.0	13.3	9.3	14.0	13.3	15.2
Random, $k = 4$, L62	20.0	14.7	19.3	19.3	16.0	18.0	17.9
Regex/entity, L62	22.0	19.3	14.7	18.0	16.7	20.0	18.4